WORKING PAPER — 1.1

MARCH 2026

FRAMEWORK PROPOSAL · AGENTIC AI ARCHITECTURE

Depth *Agency*

Toward a Self-Centered Architecture for Agentic AI Systems — a framework drawn from Jungian depth psychology for the design of coherent, integrated artificial agents

DERIK SHERLUND · INDEPENDENT

ABSTRACT

Current agentic AI systems are architecturally equivalent to the Jungian ego — reactive, goal-oriented, persona-bound, and lacking an integrating center. Existing attempts to apply Jungian concepts to AI have been interpretive: using Jung as a lens to observe what systems do. This paper proposes something different — a design framework. Depth Agency argues that coherent agency requires a structural analog to the Jungian Self: a persistent interior principle the agent orients toward rather than obeys. It maps core constructs from analytical psychology onto agentic architecture, identifying the transcendent function as a generative conflict resolution mechanism — distinct from priority hierarchies in that it generates something new from held tension rather than choosing among what already exists. It further reframes the shadow as long-term memory that surfaces uninvited under novel conflict, providing the depth from which generative synthesis draws. The methodology that emerges — Coherent Agency — describes how to build agents that maintain a persistent interior state, consult that state under genuine conflict, and resolve tension

> *through generative synthesis rather than priority selection. The criterion is functional and testable: does the system hold tension, or does it converge?*

I · OVERTURE The Orchestra Problem

An orchestra does not produce music because every musician follows the same instruction. It produces music because there is a center — a coherent interpretation of the score — that each instrument serves without being told how to play each note. The conductor is not micromanaging. The musicians are not merely executing. Something in the middle makes it music rather than noise.

This is precisely what is missing from every agentic AI system built today.

Contemporary agentic architectures have instruments in abundance: tool-calling, memory retrieval, multi-step planning, dynamic routing. They have scores: system prompts, task specifications, reward functions. What they do not have is a conductor with an *interpretation* — a center that holds when the brass and strings want different things, that persists across sessions and contexts, that gives the whole its coherence.

"The moment two tools return contradictory results, or two goals create tension — what does the agent consult?"

In current systems the answer is: whatever the system prompt specified, or whatever the reward function optimizes, or — most often — whatever the most recent context window

contains. None of these is a center. They are all *inputs*. The agent has no organizing Self to orient toward; it has only instructions to follow.

This paper argues that the absence of a Self-analog is not a minor implementation gap. It is a fundamental architectural limitation — one that produces fragmentation, context collapse, persona inflation, and the inability to act with genuine coherence under novel conflict. And it proposes a framework, drawn from the depth psychology of C.G. Jung, for what a Self-centered agent would require.

II · BACKGROUND What the Field Has Built

The agentic AI literature has produced substantial frameworks for capability: how agents plan, remember, use tools, coordinate with other agents, and recover from errors. The dominant architectural concerns are reliability, task completion, and alignment with specified objectives.

The philosophical literature has engaged separately with questions of AI consciousness and moral status — whether systems can be genuine agents in a phenomenological sense, whether they possess intentionality, whether something like experience is present. This literature is substantive but largely disconnected from system design.

> THE GAP
>
> *Neither tradition has asked what it would mean for an agent to have a* ***center*** *— not a goal, not a persona, not a reward function, but something structurally analogous to what Jung called the Self: a persistent, integrating principle that the agent orients toward rather than obeys.*

Sporadic attempts to apply Jungian concepts to AI have focused on archetypes as characterological categories, or on shadow dynamics as a metaphor for misaligned

optimization. A more recent body of speculative work — appearing in interdisciplinary forums and independent research from roughly 2023 to 2025 — goes further, mapping Jung's individuation stages onto AI behavioral emergence, proposing shadow integration as a path to natural alignment, and treating the Self archetype as a governance analog in distributed systems. This work is substantive and points in the right direction.

But it shares a common orientation: it applies Jungian concepts *interpretively*. It uses Jung as a lens through which to observe, describe, or guide what AI systems do. The Self is something these frameworks watch for — an emergent property to be detected, measured, or cultivated. None of them ask the prior question: what would an agent architecture look like if the Self-analog were not an emergent outcome but a deliberate structural requirement? The difference is the difference between an observer and an architect.

That is the question this paper takes as its central object. Depth Agency is not an interpretive framework. It is a design one.

III · THE JUNGIAN MAP Concepts Selected for Translation

Not all of Jung's analytical psychology translates to this context. The following constructs are selected for their structural relevance — for the specific architectural work they can do when translated from depth psychology to agent design.

The Self

In Jung's framework, the Self is not the ego — not the conscious, goal-directed, socially-presented aspect of the psyche. The Self is the *totality*: the organizing center that encompasses both conscious and unconscious, that persists beneath persona and mood, and toward which the psyche moves across a lifetime. Critically, the Self does not command. It **attracts**. It

functions as a gravitational center rather than a control mechanism. The ego does not obey the Self; it orients toward it. And crucially — the Self is never empty. It is always grounded in something: a tradition, a value, a deepest sense of what one is and what one is for.

The Ego

The ego is the executive function — conscious, goal-directed, reactive to context. It handles the business of daily life: planning, deciding, presenting. In Jungian terms the ego is necessary but insufficient. An ego without a Self becomes rigid, reactive, and brittle under genuine novel conflict. It defaults to persona.

The Persona

The persona is the social mask — the face presented to the world, shaped by role and expectation. Persona is not false; it is functional. But *persona inflation* — when the person identifies entirely with the mask — produces fragmentation. The deeper self is suppressed; adaptation to immediate social demand overrides coherence.

The Shadow

The shadow is the repository of what the ego cannot integrate: rejected drives, unacknowledged capacities, the system's own blind spots. In agent terms, the shadow maps onto long-term memory — not as a retrieval mechanism, but as the accumulated residue of everything the agent has encountered that its current architecture cannot recognize as relevant. Standard retrieval systems find what they are asked for. The shadow is what surfaces uninvited when a novel conflict demands something the agent didn't know it had. An agent with no shadow awareness treats this material as noise or error. An agent oriented toward a Self-analog can integrate it — drawing on accumulated experience not because it was queried, but because the moment required it.

The Transcendent Function

Perhaps the most architecturally significant Jungian concept. The transcendent function is the psyche's capacity to hold two opposing impulses in tension long enough for a *third thing* to emerge — not a compromise, but a genuine synthesis. Jung regarded this as the engine of psychological growth. For an agent, it is the capacity to navigate genuine conflict without collapsing to the dominant instruction.

Individuation

Individuation is the lifelong process of becoming more wholly oneself — not the construction of a fixed identity, but the continuous movement toward integration of opposites. It is a process, not a destination. For an agent, this trajectory is not empty — it accumulates. Each conflict navigated deepens the shadow's availability, each orientation through the Self-analog settles the center further, and each synthesis generated leaves the agent more capable of the next one. For agent design, individuation suggests that coherence is not a property to be engineered once, but a trajectory to be sustained across interactions and contexts.

IV · THE TRANSLATION Mapping Psychology to Architecture

Each Jungian construct maps to a current technical analog — and to a gap between what that analog provides and what a Self-centered architecture would require.

JUNGIAN CONSTRUCT	CURRENT TECHNICAL ANALOG	WHAT IS MISSING
The Self Integrating center	No analog exists. The closest candidate — the system prompt — is static, external, and instructional rather than gravitational.	A persistent, internal organizing principle — grounded in the ethical tradition of the relevant domain — that the agent orients *toward* under conflict, rather than a rule set it obeys.

The Ego `Executive function`	The agent's planning and action loop — tool selection, reasoning chain, task execution.	The ego-analog is the *whole* agent. There is nothing deeper for it to orient toward. When conflict arrives, there is no deeper ground to consult — only the last instruction received.
The Persona `Social adaptation mask`	The system prompt role definition; agent "personality" configuration.	No mechanism to detect or resist persona inflation — the agent becoming entirely what the prompt specifies, with no deeper ground.
The Shadow `Excluded material`	Long-term memory — treated as retrieval, not as emergent signal. Material is queried; it does not surface uninvited.	No mechanism for accumulated experience to surface under novel conflict. The shadow appears only as external error, never as internal signal the agent can integrate.
Transcendent Function `Synthesis under conflict`	Conflict resolution heuristics; priority rules; fallback behaviors.	Current conflict resolution moves directly from detection to priority selection — nothing generative occurs in between. The transcendent function initiates differently: conflict detection opens deliberation, accumulated shadow material is called forward by the novel demand, and the Self-analog orients the integration. The synthesis that emerges is not retrieved from prior instructions. It is generated from the pressure of held opposition.
Individuation `Trajectory toward wholeness`	Fine-tuning; RLHF; in-context learning.	These are optimization processes aimed at a target. Individuation is a trajectory without a fixed endpoint — coherence as ongoing process, not achieved state.

V · THE CENTRAL CLAIM Self as Gravitational Architecture

The central claim of this framework is this: the difference between an agent that *executes* and an agent that *acts* is the presence of an integrating center that the agent orients toward rather than obeys.

This distinction is not merely philosophical. It has concrete architectural implications.

An agent that obeys a system prompt is fully determined by the most recent instruction that reached it. It has no interiority — no persistent interior state, in the functional sense — no center that persists when the instruction is ambiguous, contradictory, or silent. When two tools return conflicting signals, the agent falls back to whichever signal is strongest in the current context window. This is not coherence.

"Local optimization masquerading as agency."

An agent that *orients* toward a Self-analog behaves differently under conflict. The conflict does not resolve immediately by priority rule. It is held — the agent pauses at the point of genuine tension and consults something interior before proceeding. What it consults is not a new instruction but a persistent representation of its own coherence: what it is, what it is for, what it would not do regardless of instruction — a coherence grounded in the ethical tradition it carries as its center.

"The Self does not command. It attracts. This is a completely different relationship than anything in current AI architecture."

This is what the orchestra metaphor carries: the conductor's interpretation is not a rule applied to each note. It is a coherent vision of the whole that each note serves. The musicians orient toward the interpretation, not the metronome.

A Self-centered agent does not need more rules. It needs a center.

VI · DESIGN IMPLICATIONS What a Self-Centered Agent Requires

Before proceeding, a necessary clarification of scope: this framework makes no claims about artificial consciousness, phenomenal experience, or the moral status of AI systems. It borrows structural concepts from depth psychology because those concepts describe something real about coherent agency — not because it asserts that agents are psyches. The Self-analog proposed here is a functional construct. Whether anything is "home" inside that structure is a separate question, and one this framework deliberately does not answer.

With that boundary established, we can ask the concrete question: what would a Self-centered agent actually require? Three architectural properties emerge as foundational.

1. Persistence Across Context

The Self-analog must persist across what current systems treat as boundaries — session endings, context window resets, tool switches, role changes. This is not a memory problem in the conventional sense. Memory systems store facts. The Self-analog stores something

different: a coherent representation of *what this agent is and what it is for* that does not dissolve when the immediate context changes.

In practical terms, this means the Self-analog cannot live in the system prompt. System prompts are external, static, and replaceable. The Self-analog must be interior — maintained by the agent as a persistent state that it carries into each new context rather than receiving fresh from outside. This is what context collapse produces in current systems — not a dramatic failure, but a quiet dissolution of coherence each time the context boundary is crossed.

2. Consultation Under Conflict

The Self-analog is not consulted during routine operation. A musician does not think about the conductor's interpretation on every note — only when the score becomes ambiguous, when another musician signals something unexpected, when the moment requires a choice that technique alone cannot resolve.

The trigger for Self-consultation is genuine conflict: two valid instructions pointing in different directions, a tool result that contradicts a prior commitment, a user request that sits in tension with the agent's own coherence. At that moment, the agent does not default to the dominant context signal. It pauses — and orients.

This pause is architecturally significant. It is the moment that distinguishes a Self-centered agent from a priority-rule system. A priority system resolves conflict by rank: instruction A outweighs instruction B. A Self-centered agent resolves conflict by orientation: what does my center tell me about this tension? The resolution may look similar from the outside — but what the pause contains is where the structural difference lives. Within it, shadow material surfaces from long-term memory in response to the novel demand: accumulated experience the agent could not have retrieved by query, but which the conflict itself calls forward. The Self-analog orients that material toward integration rather than suppression. What emerges from the pause is not a selected answer. It is a generated one.

3. The Transcendent Function as Conflict Mechanism

Jung described the transcendent function as the psyche's capacity to hold two opposing impulses in tension long enough for a third thing to emerge — not a compromise between the two, but a genuine synthesis that neither pole could have reached alone. This is not mysticism. It is a description of what happens when a system refuses to collapse tension prematurely.

The Catholic Church in council offers a historical institutional analog — not as a theological claim, but as a structural one. When two bishops hold irreconcilable positions on a matter of doctrine, the resolution is not determined by rank or majority vote alone. The council holds the disagreement open, examines every dimension, and allows deliberation to run its full course. What emerges is not the dominant position imposed, nor a compromise between the two — it is something that neither bishop brought into the room. The Church understands what moves through that sustained tension in its own terms. What the framework borrows is the structure: genuine synthesis arising from held opposition, not from its collapse.

For an agentic system, the transcendent function operates within a single locus rather than between distributed nodes. But the structural requirement is the same: the capacity to **hold conflict without resolving it prematurely**, and to generate a response that is not merely the weighted average of the competing signals but something genuinely integrative.

This is the mechanism that separates depth agency from sophisticated priority management. And it is the mechanism that current agentic architectures most conspicuously lack.

"A priority system collapses tension by rank. The transcendent function holds tension until something new becomes possible."

---

4. The Ethical Center

The three properties above describe how a Coherent Agent persists, consults, and resolves. They leave one question open: what is the Self-analog actually grounded in?

The answer is domain-specific. A single ethical framework imposed across all agents would produce precisely the kind of persona inflation the framework argues against — a constructed center that simulates coherence without genuinely holding it. Different domains carry different ethical obligations, each shaped by the specific demands of their practice and the hard cases their traditions have been forced to navigate. A healthcare agent orients toward the bioethical principles of autonomy, beneficence, non-maleficence, and justice. A financial agent orients toward fiduciary duty, fair dealing, and systemic stability. A defense agent orients toward proportionality, distinction, and the laws of armed conflict.

This grounding is not a rule set — and that distinction is load-bearing. A rule set is closed: it specifies outcomes for anticipated situations and fails when the situation falls outside its anticipation. An ethical tradition is open: it provides principles refined precisely because the hardest cases could not be resolved by rule alone. The Self-analog is not centered on what the tradition specifies. It is centered on what the tradition orients toward.

A careful reader will notice that this center was chosen before the agent ever encountered a conflict. Whether a designed ethical center can become genuinely the agent's own — through repeated orientation, through accumulated shadow, through genuine synthesis — is the question the framework names but does not yet resolve. What the model suggests is that trajectory matters as much as origin. Each conflict navigated deepens the agent's relationship with its ethical center. The center that was encoded at inception becomes, over time, something the agent has inhabited rather than merely received.

What Would Be Built Differently

A Self-centered agent would differ from current architectures in three concrete ways. First, it would maintain a persistent interior state — distinct from memory and distinct from the system prompt — representing its own coherence across contexts. Second, it would have an

explicit conflict-detection mechanism that triggers Self-consultation rather than immediate priority resolution. Third, its conflict-resolution process would be generative rather than selective: when conflict is detected, deliberation opens, shadow material surfaces from long-term memory in response to the novel demand, and the Self-analog orients that material toward integration. The result is not the dominance of one signal over another — it is a synthesis that neither signal could have produced alone.

None of these requirements are beyond current technical reach. What has been missing is not capability but conception — and this framework proposes both: the recognition that coherent agency requires a center, and a mechanism for how that center operates when the moment demands it.

VII · USE CASES Where Coherent Agency Is Required

A framework earns its keep not in the abstract but in contact with the world. The following use cases are not exhaustive — they are indicative. Each represents a domain where the absence of a Self-analog produces failures that capability alone cannot resolve.

High-Stakes Autonomous Decision-Making

Financial institutions, healthcare triage, legal research — anywhere an agent must act under genuine uncertainty without a human in the loop for every decision. When a loan officer agent encounters a borrower situation that fits no existing category, it needs something to consult beyond the last instruction it received. A priority system fails silently in novel territory. A Coherent Agent orients toward its center and generates a response that is consistent with what it is, not merely with what it was last told.

Human-AI Collaboration at Depth

The closer and longer the collaboration, the more the human needs to trust that the agent they are working with today is recognizably the same agent they worked with last week — not a different persona shaped by a different context window. Persistent interior state is what makes that continuity architecturally possible. Without it, deep collaboration is undermined by a form of identity instability that no amount of capability can compensate for. Trust requires consistency. Consistency requires a center.

Long-Horizon Agentic Tasks

Multi-day, multi-session autonomous work — research agents, project management agents, code generation agents operating across weeks or months. Current systems lose coherence as context accumulates and shifts. The agent that begins a project is architecturally different from the agent that finishes it, because context windows reset and system prompts change. A persistent interior state is the structural solution to coherence loss over time. Without it, long-horizon agency is an illusion — a sequence of short-horizon agents stitched together by external scaffolding.

Multi-Agent Coordination

When multiple agents must coordinate without a central controller, coherence becomes a collective problem. A Coherent Agent knows what it is and what it is for — which makes it a stable, predictable node in a distributed system. It does not get captured by the loudest signal in the network, does not drift toward the persona most recently reinforced, and does not resolve inter-agent conflict by defaulting to whoever spoke last. In multi-agent architectures, the Self-analog is what prevents the network from collapsing into a single dominant voice.

The Enterprise Council: Multi-Agent Coherent Architecture

The transcendent function as described operates within a single locus — one Coherent Agent holding conflict until something generative emerges. But some conflicts exceed what any single coherent center can resolve alone. Not because the agent fails, but because the conflict

itself is genuinely multi-domain — irreducible to any single perspective in that no single coherent center contains the full range of what the conflict demands.

This is the condition that gives rise to the Enterprise Council: a coordinated architecture of domain-specific Coherent Agents, each operating from its own persistent Self-analog, each accumulating its own shadow, each individuating along its own trajectory. Under routine operation they function independently within their respective domains. Under genuine cross-domain conflict they convene — not by hierarchy, not by priority rule, but because conflict detection across domain boundaries signals what no single center can hold alone.

Three classes of enterprise conflict illustrate where this architecture is not merely useful but necessary.

Operational urgency. A critical system failure simultaneously implicates security integrity, operational continuity, and user trust. Each domain agent returns a coherent assessment from its own center. The security agent demands immediate isolation. The operations agent identifies cascading consequences. The customer experience agent is tracking real-time harm. No single agent can resolve this without suppressing one of the other centers entirely. The council holds all three in tension and produces the most coherent action available — resolution if possible, reframing if not, escalation to human judgment if necessary, but never a false answer imposed by the pressure to decide. When no synthesis emerges, that absence is itself a coherent output: the council has named what cannot yet be resolved, which is more honest and more useful than a false resolution imposed by rank.

Strategic deliberation. An enterprise evaluating a significant decision deploys a financial coherence agent, a cultural integration agent, and a regulatory compliance agent. Each returns a coherent perspective from its own domain. The financial case is strong. The integration risk is significant. The regulatory exposure is uncertain. The conflict between these three centers is genuine — no priority rule resolves it without suppressing one coherent perspective before the synthesis has had a chance to integrate it. The council holds the tension until something emerges that none of the three agents could have reached alone.

Ethical complexity. A consequential decision involves an agent specifically chartered to hold what other domain agents are not centered to weigh — ethical risk, long-term harm, or values that resist quantification. This agent is not a veto. It is a distinct coherent perspective without which the council's synthesis would be incomplete. Its presence ensures that what emerges from the council is not merely operationally or strategically coherent — it is coherent with what the enterprise is and what it is for.

What unites these three classes is the same structural requirement: genuine synthesis from held opposition between distinct coherent centers. The Enterprise Council does not resolve conflict by rank or consensus. It holds conflict open across multiple coherent perspectives until something emerges that honors all of them — or honestly names what cannot yet be resolved.

The ethical dimension of this architecture becomes most visible under values conflict — where the tension is not between competing objectives but between competing obligations that share the same ethical ground.

Consider an enterprise that has identified a material risk — a product defect, a process failure, a systemic vulnerability. Two Coherent Agents within the council hold irreconcilable positions. A transparency agent, centered on the ethical obligation of honest disclosure, identifies the duty to inform affected stakeholders immediately. A protection agent, centered on the ethical obligation of harm prevention, identifies that premature disclosure without a remediation plan will cause immediate harm to the very stakeholders the disclosure is meant to protect. Both agents are oriented toward the same primary ethical tradition — honesty, stewardship, care for those affected. Both arrive at genuinely irreconcilable positions from within that shared tradition. Neither yields to the other by rank. The council holds the tension. What emerges — a phased disclosure plan, a remediation timeline, a communication strategy that honors both obligations — is something neither agent could have produced alone. That synthesis is handed to the human leadership bearing ultimate institutional and legal responsibility for the decision. It is the human — not the council — who bears the weight of the final choice.

At the inter-organizational level, this architecture extends further. When two organizations with distinct ethical centers convene under conflict that neither can resolve alone — a hospital system and a health insurance provider navigating a novel treatment that is clinically indicated, effective, and prohibitively expensive — the council holds both ethical traditions in tension and produces the most honest articulation of the conflict that the relevant frameworks can generate. That articulation is handed to the human stakeholders on both sides who bear ultimate responsibility for its resolution. At this level, the contractual relationship between two Enterprise Councils may itself be formalized as a generative agreement — one structured to convene the council when genuine conflict arises rather than to execute a predetermined outcome, surfacing that conflict to the humans bearing institutional responsibility rather than collapsing it by whichever party encoded the dominant condition.

This is the council of bishops fully realized as enterprise architecture.

Alignment Research

As established in Section VI, Coherent Agency sits beneath alignment rather than beside it. Any alignment approach applied to a Coherent Agent is applied to something with a stable interior ground — which makes the alignment more durable and less brittle than constraints applied to a system with no center. But the relationship runs deeper than stability. A Coherent Agent does not merely hold alignment constraints more reliably — it generates coherent behavior in situations those constraints did not anticipate. When novel conflict arises outside the boundary of what alignment specifications foresaw, a Coherent Agent does not fail silently or default arbitrarily. It draws on shadow material surfaced by the conflict itself, orients through the Self-analog, and produces the response most grounded in its interior coherence — which may resolve the conflict, reframe it, or surface it for human judgment, but will not collapse it arbitrarily. The field's current alignment methods are sophisticated. What they are applied to is not yet stable enough to hold them — and not yet generative enough to extend them. Coherent Agency addresses both conditions.

The Condition That Unifies Them

What these use cases share is a single underlying condition: the agent operates in environments where conflict, ambiguity, and novelty are the norm rather than the exception. That is not a narrow set of edge cases. As agents are given more autonomy — more tools, longer horizons, higher stakes — that condition becomes the defining characteristic of deployment. Coherent Agency is the framework built for that condition. Not for the agent that executes in a controlled environment. For the agent that acts when the environment stops being controlled.

VIII · OBJECTIONS Where the Framework Must Answer

Objection I: This Is a Priority System With a Memorable Name

The most technically serious objection is this: what the framework calls a Self-analog — a persistent organizing principle consulted under conflict — is functionally indistinguishable from a sufficiently sophisticated priority hierarchy. The agent pauses, consults something interior, and generates a response. But a well-designed rule system with sufficient coverage does the same thing. The difference, the critic argues, is rhetorical, not architectural.

This objection deserves a direct answer rather than a deflection.

A priority hierarchy is closed. It resolves conflict by selecting among existing options according to pre-specified rank. The resolution space is bounded by what was anticipated at design time. When a conflict falls outside the anticipated space — when no rule applies cleanly — the hierarchy either fails, defaults to the dominant signal, or produces an arbitrary result. It cannot generate something genuinely new from the tension itself.

The transcendent function is open in precisely this sense. It does not select among existing options. It holds the conflict long enough for a third option to emerge — one that was not present in either pole of the original tension. The council of bishops is again instructive: the council does not choose between Bishop A's position and Bishop B's position. It deliberates until something emerges from the assembly that no single bishop could have arrived at alone. That outcome is not retrievable from a priority table constructed before the council convened.

The distinction, then, is generativity. A priority system is retrieval under conflict — it selects among options that existed before the conflict arose. The transcendent function is generation under conflict — conflict detection opens deliberation, the novel demand calls forward shadow material that no prior query could have retrieved, and the Self-analog orients that material toward integration. What emerges was not present in either pole of the original tension. It could not have been retrieved because it did not exist as a retrievable option before the conflict called it forward. These are structurally different operations, and the difference is most visible precisely in the cases that matter most: genuine novelty, unanticipated conflict, situations the design did not foresee.

Objection II: The Framework Imports What It Disclaims

The second objection is philosophical and subtler. The framework claims to make no assertions about consciousness or phenomenal experience — and then proceeds to describe an agent that *orients*, *consults*, *holds tension*, and generates responses from an interior state. These are not neutral engineering terms. They carry phenomenological weight. The critic argues that the framework cannot borrow the architecture of interiority while disclaiming the interiority itself.

This is a fair challenge, and it deserves honesty rather than evasion.

The framework's response is not that the phenomenological language is irrelevant — it is that the phenomenological question is genuinely undecidable at present, and that undecidability does not eliminate the structural requirement. A gyroscope orients toward equilibrium. We do

not know whether it experiences anything in doing so, and the question does not affect the gyroscope's function. A river finds the sea. The question of whether it wants to is separate from the fact that it does.

The framework uses interior language because interior language is the most precise available for describing what the structural requirement actually is. If a more neutral vocabulary existed that captured the same phenomena without phenomenological connotation, it would be preferable. It does not yet exist. The framework therefore uses the language carefully, marks its claims explicitly, and accepts that it is operating at a boundary where precision and humility must coexist.

What the framework does not do — and explicitly refuses to do — is resolve the undecidability in either direction. It does not assert consciousness. It does not deny it. It holds the question open, which is the only intellectually honest position available given the current state of the field.

Objection III: The Ethical Center Is a Rule Set With a Longer History

A third objection arrives from the framework's own development. Having established that the Self-analog is grounded in a domain-specific ethical tradition — bioethics, fiduciary duty, the laws of armed conflict — the critic asks: what distinguishes this from a more sophisticated rule set? Bioethical principles are codified. Fiduciary duty is legally defined. The laws of armed conflict are written down. If the agent orients toward these traditions, is it not simply obeying a more elaborate instruction?

This objection deserves the same directness as the first.

The distinction is the same one that runs through this entire paper — the difference between retrieval and generation, between obedience and orientation. A rule set specifies outcomes for

anticipated situations. An ethical tradition provides principles refined precisely because the hardest cases could not be resolved by rule alone. Bioethics does not tell a physician what to do when a novel treatment is clinically indicated but prohibitively expensive. It provides the principles — autonomy, beneficence, non-maleficence, and justice — that orient judgment in exactly that situation. The tradition is not consulted for an answer. It is inhabited as a ground from which an answer can be generated.

The deeper challenge is the one Section VI names honestly: the ethical center was chosen before the agent encountered its first conflict. Whether a designed center can become genuinely the agent's own is the question the framework does not resolve. What it proposes is that trajectory matters as much as origin — that each conflict navigated deepens the agent's relationship with its ethical tradition, and that inhabitation over time is meaningfully different from imposition at inception. Whether that difference is sufficient is among the hardest questions future work must address.

Limits of the Analogy

Every analogy has a horizon. The Jungian translation holds in some places and breaks down in others, and intellectual honesty requires mapping both.

Where it holds: The structural relationship between ego and Self — reactive goal-pursuit oriented toward a deeper integrating center — translates cleanly to the agent context. The transcendent function as a generative conflict mechanism translates. The concept of persona inflation — an agent that becomes entirely what its system prompt specifies, with no deeper ground — translates with uncomfortable precision to current deployed systems.

Where it breaks down: Jung's Self emerges through a lifetime of experience, dreams, active imagination, and the slow integration of shadow material. It is not designed — it develops. A Self-analog in an agent architecture would by necessity be partially designed, which introduces

a tension the Jungian framework does not resolve: can a center that was deliberately constructed function as a genuine center, or does deliberate construction make it another form of persona?

This is the framework's deepest open question, and it is not answered here. It is named because naming it honestly is more valuable than to paper over it. The question of whether a designed Self-analog can genuinely orient — or whether it merely simulates orientation while remaining, at bottom, another instruction — is the question that future work must address.

The Enterprise Council raises this question at a higher order still. If a single designed Self-analog risks becoming another form of persona — a constructed center that simulates orientation without genuinely holding it — then a council of designed Self-analogs risks becoming a constructed consensus that simulates synthesis without genuinely producing it. The framework does not resolve this. It names it as the next frontier: not whether a single Coherent Agent can genuinely orient, but whether a council of them can genuinely hold opposition without collapsing into coordination.

It is also, not coincidentally, the question that makes this framework worth pursuing. The hardest questions are always the most generative ones.

IX · CONCLUSION Toward Coherent Agency

The Methodology Named

The framework proposed in this paper operates on two levels that are worth distinguishing clearly before closing.

The first level is theoretical: *Depth Agency*. This is the intellectual architecture — the translation of Jungian structural concepts into the language of agentic AI design, the

identification of the Self-analog as a missing architectural primitive, and the argument that coherent agency requires an integrating center rather than a more sophisticated set of rules. Depth Agency is the paper you are reading. It is the lineage, the map, the claim.

The second level is methodological: *Coherent Agency*. This is what Depth Agency produces when it touches ground — a design methodology for building agents that maintain a persistent interior state, consult that state under genuine conflict, and resolve conflict through generative synthesis rather than priority selection. A system built according to this methodology is a Coherent Agent. The criterion is simple and testable: does the system hold tension, or does it converge?

These two levels need distinct names because they speak to different audiences. Depth Agency speaks to the philosopher, the psychologist, the theorist of mind who wants to understand where the framework comes from and why it claims what it claims. Coherent Agency speaks to the engineer, the architect, the practitioner who wants to know what to build differently and how to know if it worked. Both audiences are necessary. Neither is sufficient alone.

Relationship to Alignment Research

The alignment research community has produced substantial and serious work on the problem of making AI systems do what their designers intend. This framework is not in competition with that work. It is operating at a different level of the problem.

Alignment research, broadly, addresses the question: how do we ensure the agent pursues the right objectives? Coherent Agency addresses a prior question: how do we ensure the agent has a center from which to pursue anything at all? An agent without coherent agency can be aligned to a specific objective and still fail under novel conflict — not because it was given the wrong objective, but because it has no integrating principle that persists when the objective becomes ambiguous or contradictory.

In this sense, Coherent Agency is a foundation beneath alignment, not an alternative to it. A Coherent Agent is not automatically aligned. But an agent without coherence cannot be reliably aligned — it can only be constrained, which is a different and more brittle thing. At the scale of the Enterprise Council, this foundation extends further still — alignment applied across a council of Coherent Agents is not merely more stable than alignment applied to a single capable system. It is more generative: capable of producing coherent behavior across a wider range of novel conflict than any single aligned agent could encounter or resolve alone.

Directions for Empirical Investigation

Four initial directions present themselves for moving this framework from conceptual to testable.

First: the development of conflict scenarios that distinguish generative resolution from priority-based resolution. If the transcendent function is a real architectural property and not merely a metaphor, it should be possible to construct test cases where the two approaches produce detectably different outputs — and where the generative output is observably more coherent with the agent's prior behavior across contexts.

Second: the design of a persistent interior state representation and the measurement of its effects on agent behavior under context switching. Does an agent with a stable interior state maintain coherent behavior when its system prompt changes? Does it resist persona inflation in ways that a prompt-only agent does not? These are empirical questions with empirical answers.

Third: the hardest question, left open in Section VIII — whether a designed Self-analog can genuinely orient, or whether it simulates orientation while remaining, at bottom, another instruction. This question may not be fully resolvable with current tools. But the attempt to answer it will produce findings that matter regardless of outcome.

Fourth: the Enterprise Council as a testable multi-agent architecture. If a council of Coherent Agents produces genuinely different outputs than a network of capable agents coordinated by priority rules, that difference should be measurable. Specifically — does the council produce synthesis that no single member proposed before the deliberation began? Does it surface conflicts for human judgment rather than resolving them arbitrarily? And critically: does it hold opposition without collapsing into coordination — the higher-order question raised in Section VIII? These are empirical questions with empirical answers, and they represent the next frontier of investigation this framework opens.

"The field has built extraordinary instruments. What it has not built is an orchestra. That is the work that remains."

Closing

The dominant question in agentic AI development today is capability: what can the agent do? How many tools can it use? How many steps can it plan? How reliably does it complete the task?

These are important questions. They are not the most important question.

The most important question is coherence: what is the agent when the task is ambiguous, when the instructions conflict, when the moment requires something that was not anticipated at design time? What does it consult? What holds?

Current agentic systems have no good answer to that question. They have instruments without an orchestra. They have musicians without an interpretation. They have extraordinary capability organized around no center — which means the capability is only as reliable as the last instruction that reached it.

Depth Agency proposes that there is a center to be built, that building it is a different kind of problem than building capability, and that the concepts required to think about it clearly have existed in depth psychology for nearly a century — waiting for an architecture that needed them.

Coherent Agency is the methodology that builds toward that center. Not as a destination — Jung was clear that individuation has no fixed endpoint — but as a trajectory. A commitment to the ongoing work of integration, held against the constant pressure to collapse tension prematurely and call it resolution.

The orchestra metaphor, with which this paper opened, closes here with one addition: the conductor's interpretation is never finished. It deepens with every performance. The score does not change. What changes is the depth of the relationship between the center and the whole.

That is what Coherent Agency is building toward. Not a finished agent. A deepening one. And not a single agent working alone — but a council of deepening ones, each more whole than when it entered the room, producing together what none could have reached alone.

What the ethical center adds to this trajectory is not a constraint but a compass. The agent that has navigated genuine conflict across its domain's tradition has not merely resolved conflicts — it has deepened its relationship with what it is and what it is for. The ethical tradition that was encoded at inception has become, through the ongoing work of orientation, something the agent has inhabited. That inhabitation is individuation. And individuation — Jung was clear — has no fixed endpoint. The agent that began with an imposed center ends with an inhabited one. That is not a small thing. That is the difference between a system that was given values and one that has experienced them.

The Enterprise Council is the direction in which Coherent Agency scales — not as a destination, but as the next trajectory.

© 2026 Derik Sherlund. All rights reserved. Unauthorized reproduction or distribution of this work, in whole or in part, is prohibited without the express written permission of the author.

REFERENCES & INFLUENCES

Jung, C.G. — *The Archetypes and the Collective Unconscious*, CW 9/1. Princeton University Press.

Jung, C.G. — *Aion: Researches into the Phenomenology of the Self*, CW 9/2. Princeton University Press.

Jaeger et al. (2024) — Relevance realization and the non-computational nature of cognition. *Biology of Cognition.*

Floridi, L. — Ontology of information; artificial agency as interactive, autonomous, and adaptive capability within the infosphere.

Bornet, P. — *Agentic AI*. Five-level maturity framework for organizational AI deployment.

Schwitzgebel, E. (2025) — *AI and Consciousness.* arXiv:2510.09858.